

Fair Ranking in Retrieval-Augmented Generation: Is the Benefit Fairness, or Diversity in Disguise?

Idan Peretz

Avi Simkin

Abstract

Kim and Diaz (2025) report that Retrieval-Augmented Generation (RAG) systems built on item-exposure-fair rankings can preserve, and in many cases even improve, downstream generation quality relative to relevance-only rankings. We investigate a specific mechanistic ambiguity in that claim: the stochastic fairness mechanism Kim and Diaz use for this purpose, Plackett–Luce (PL) reranking, does not only redistribute expected exposure across repeated rankings — as a byproduct of its randomization, it also tends to change the within-list textual diversity of what a RAG generator actually sees. If diversity itself helps generation, an apparent “fairness helps” result could really be diversity doing the work. We test this directly on the LaMP benchmark (all 7 tasks, both generators Flan-T5-Small and Flan-T5-Base, both BM25 and Contriever retrieval, all available queries retained by Kim and Diaz’s generator-specific fairness-evaluation collections for these two generators, 336 experiment configurations, 163,392 pre-exclusion query-setting rows) by using Maximal Marginal Relevance (MMR) — a reranker with an explicit diversity dial and no fairness objective — to separate the two constructs. We find no general or consistent utility improvement from PL reranking over the deterministic baseline (10 of 16 Holm-corrected tests are significant: 8 of those 10 are losses and only 2 small gains); that within PL’s own fairness dial, greater exposure equality is robustly associated with *lower* utility (coefficient $+0.0735$, cluster-robust $p = 6.3 \times 10^{-54}$, stable under the pooled specification and two progressively stricter fixed-effects specifications); that MMR-induced diversity shows no general utility benefit of its own (a modest, heterogeneous negative pooled association, $p = 0.0115$); and that PL retains an average utility disadvantage relative to MMR after adjusting for measured diversity (average marginal effect ≈ -0.019 to -0.020 , robust across every specification and every leave-one-task-out fold we tested). A directly tested candidate mechanism — that PL pays a relevance cost via reduced exposure-relevance alignment — does not explain this residual. We conclude that, in the settings we test, measured within-list diversity (ILD-Jaccard) alone does not account for the difference between PL and MMR utility, while explicitly declining to interpret this residual as the causal cost of fairness itself, since the two rerankers differ along dimensions besides diversity.

1 Introduction

Retrieval-Augmented Generation (RAG) systems answer a query by first retrieving a small set of supporting documents or profile items and then conditioning a language model’s generation on that retrieved set (Lewis et al., 2020). Because only a handful of items are shown to the generator, *which* items are retrieved and in *what order* they are presented has a direct effect on the generator’s output. This same property has made ranking a natural place to intervene for *item-side fairness*: if being retrieved and shown carries value (exposure) to the entity that produced an item, then systematically favoring some relevant items over other, comparably relevant items is a fairness concern independent of whether it also affects end-task quality. A substantial ranking-fairness literature formalizes this as exposure fairness — equalizing the *expected exposure* an item receives

across repeated, possibly randomized rankings, rather than requiring a single fixed ranking to look fair on its own (Singh and Joachims, 2018; Biega et al., 2018; Diaz et al., 2020).

Kim and Diaz (2025) apply this idea to RAG directly, reranking retrieved candidate sets with fairness-aware methods before generation and measuring both item-side exposure fairness and downstream generation quality. Their central, and somewhat counter-intuitive, finding is that RAG systems built on fair rankings can preserve, and in many cases even outperform, RAG systems built on purely relevance-ranked retrieval — suggesting that enforcing fairness need not cost generation quality, and might sometimes help it.

This project investigates a specific reason that result could be misleading. The stochastic fairness mechanism Kim and Diaz use, Plackett–Luce (PL) reranking, redistributes expected exposure across repeated draws by randomizing which comparably relevant items are shown, which can reduce exposure disparity as the sampling distribution flattens (Plackett, 1975; Luce, 1959). That randomization is not targeted at diversity in any direct sense, but as a side effect it tends to change how varied the items shown *in any single sampled list* are from one another. If a RAG generator produces better output when it sees a more varied context — independent of whether that context is fairly exposed — then an observed “fair reranking helped” result could really be within-list diversity doing the work, with fairness merely along for the ride. Before attributing a utility change to fairness, it is therefore necessary to check whether the diversity that fairness happens to induce is the actual mechanism.

This is the question this paper answers:

When a fair ranking method appears to help, or hurt, a RAG system’s output quality, is that effect attributable to fairness itself, or to the within-list diversity that the fairness mechanism happens to produce?

Answering this requires separating two constructs that are easy to conflate in practice but are conceptually distinct: *fairness*, which we operationalize as expected exposure equality across repeated rankings (Expected Exposure Disparity, EE-D, lower is fairer), and *diversity*, which we operationalize as within-list textual variety in a single ranking (ILD-Jaccard). PL moves both at once. MMR provides a complementary diversity-focused comparator (Carbonell and Goldstein, 1998): it has an explicit relevance–diversity dial and no exposure–fairness objective. Comparing PL and MMR therefore helps test whether measured within-list diversity is sufficient to account for PL’s utility pattern, while not causally isolating diversity because changing MMR’s λ also changes the relevance–diversity tradeoff — though, as we discuss throughout, any residual difference between PL and MMR is not automatically identifiable as “the causal effect of fairness,” since PL and MMR differ along several dimensions besides diversity (stochastic versus deterministic selection, how each preserves the underlying relevance ordering, and so on).

Contributions. We (1) reproduce and re-examine Kim and Diaz’s central claim on a subset of their experimental design (all 7 LaMP tasks, two generators, two retrievers, using all available queries retained by their generator-specific fairness-evaluation collections for these two generators); (2) test, rather than assume, whether PL’s fairness mechanism actually induces the diversity increase the diversity-as-mechanism hypothesis requires, at both the aggregate and per-query level; (3) use MMR as a diversity-focused comparator with an explicit diversity dial and no fairness objective to test whether measured within-list diversity is sufficient to account for PL’s utility behavior, reporting four main findings with cluster-robust inference, leave-one-task-out robustness, and two progressively stricter levels of fixed-effects identification on top of the pooled specification; and (4) directly test our leading hypothesis for the residual PL–MMR utility gap (a relevance cost from

PL’s stochastic sampling) and find no support for it, leaving the mechanism explicitly open rather than overclaiming an explanation.

2 Background and Related Work

Retrieval-augmented generation. RAG conditions a generation model on a small set of retrieved documents or profile items rather than relying solely on parametric knowledge (Lewis et al., 2020). This makes generation quality directly sensitive to the retrieval and reranking pipeline: which items are retrieved, and the order and composition of the final shortlist handed to the generator, can change what the model is able to produce. Reranking the initial candidate set — the stage this paper studies — is therefore a natural place to intervene on both fairness and diversity objectives without retraining the generator itself.

Exposure fairness in ranking. A ranking necessarily privileges higher-ranked items with more user attention; item-side fairness concerns whether comparably relevant items receive comparable exposure over repeated exposure to the ranking, not whether any single ranking looks balanced in isolation (Singh and Joachims, 2018; Biega et al., 2018). Diaz et al. (2020) formalize this with the *expected exposure* framework: each (possibly randomly sampled) ranked list induces a *system exposure vector*, one entry per candidate item, and item relevance induces an analogous *target exposure vector*. Under this framework’s uniform top- k exposure model — the one this paper’s implementation uses throughout — every item placed in the top k positions of a sampled ranking receives exposure 1 and every item outside the top k receives exposure 0, averaged across however many rankings are sampled for a query (a single ranking for the deterministic baseline and MMR; the full set of sampled rankings for PL, Section 4). This paper’s fairness metric, EE-D (Expected Exposure Disparity), is the squared norm of the system exposure vector itself: it is higher when exposure concentrates repeatedly on the same few items and lower when it spreads across many different items, independent of which items are relevant. EE-R (Expected Exposure Relevance, used only in the mechanism test of Section 5.4) is the inner product between the system and target exposure vectors: it is higher when exposure is better aligned with useful items. EE-D is therefore not itself a distance between delivered and target exposure — that quantity is a separate metric this paper does not use. Both EE-D and EE-R are computed with the reference implementation released alongside the expected-exposure framework, applied unmodified to this project’s rankings.

Kim and Diaz’s Fair RAG. Kim and Diaz (2025) evaluate twelve RAG models (3 retrievers — BM25, SPLADE, and Contriever — crossed with 4 generators — Flan-T5-Small, Flan-T5-Base, Flan-T5-XXL, and Flan-UL2) across the fairness metrics above and downstream task utility, using the LaMP personalization benchmark (Salemi et al., 2024) augmented with their own item-level utility labels. Their central finding — that fair rerankings can preserve or even improve generation quality relative to purely relevance-ranked retrieval — motivates the present paper’s question: whether that apparent quality preservation is attributable to fairness itself, or to a byproduct of the fairness mechanism used to achieve it.

Plackett–Luce stochastic ranking. The Plackett–Luce (PL) model (Plackett, 1975; Luce, 1959) defines a distribution over permutations of a candidate set by sampling without replacement. Following Kim and Diaz’s pipeline, raw retrieval scores for a query are first min–max normalized to $[1, 2]$ (a special case rescales binary oracle relevance instead, not used for the BM25/Contriever rankings in this paper), giving a normalized score $\bar{s}_d$ for each candidate item d . A *fairness-control*

parameter α is then applied to each normalized score by exponentiation, $\bar{s}_d^\alpha$, before sequential Plackett–Luce sampling (implemented via the Gumbel-max trick): at each step, an item d is selected from the remaining candidates R with probability

$$p(d \mid R) \propto \exp(\bar{s}_d^\alpha).$$

Because $\bar{s}_d \in [1, 2]$, a larger α stretches the gap between higher- and lower-scored items’ transformed scores, sharpening the sampling distribution toward the highest-scored item and producing more deterministic, higher-disparity rankings — a weaker fairness intervention. A smaller α flattens the transformed scores toward each other, producing lower disparity and a stronger fairness intervention. We avoid describing α as a conventional softmax temperature, since it is not applied as a divisor of the scores before exponentiation in the usual temperature-scaling sense; it is a power applied to the normalized score itself. Repeated draws from this distribution, rather than a single deterministic ranking, are what makes expected-exposure equality achievable: items with comparable transformed scores receive comparable exposure *in expectation* even though any single sampled ranking still favors higher-scored items.

Maximal Marginal Relevance. MMR (Carbonell and Goldstein, 1998) constructs a ranking greedily: having selected items $D_1, \dots, D_{k-1}$, the next item is

$$D_k = \arg \max_{D_i \notin \{D_1, \dots, D_{k-1}\}} \left[\lambda \cdot \text{Sim}_1(D_i, Q) - (1 - \lambda) \cdot \max_{D_j \in \{D_1, \dots, D_{k-1}\}} \text{Sim}_2(D_i, D_j) \right],$$

where Sim_1 is relevance to the query and Sim_2 is similarity to already-selected items. In our implementation, Sim_1 is the raw retrieval score (BM25 or Contriever), min–max normalized to $[0, 1]$ per query — a separate, local normalization from the $[1, 2]$ rescaling PL’s sampler uses above. Sim_2 is the Jaccard similarity between two items’ token sets, where each item’s token set is every string-valued profile field except its identifier and date fields, lowercased and split on whitespace (no stemming or stopword removal); this is the identical tokenization used for the ILD-Jaccard diversity metric itself (Section 3). The parameter $\lambda \in [0, 1]$ trades relevance for diversity: $\lambda = 1$ recovers the purely relevance-ranked list, and decreasing λ pushes the ranking toward more mutually dissimilar items. MMR is entirely deterministic and has no exposure-fairness objective, which makes it a useful diversity-focused comparator for testing whether measured within-list diversity alone is sufficient to account for PL’s utility behavior; because changing λ also changes the relevance-diversity tradeoff, however, MMR does not causally isolate diversity by itself.

3 Methodology

3.1 Datasets

We use the LaMP personalization benchmark (Salemi et al., 2024), all seven tasks, under the “user” split, drawing all available queries per task from the generator-dependent, per-generator-filtered utility-label test collections that Kim and Diaz’s released data pipeline builds from the original LaMP release (Kim and Diaz, 2025); we do not draw any additional subsample on top of that filtering.

Kim and Diaz (2025) construct these collections as follows. Item-level usefulness labels are defined by marginal generator utility: each profile item is individually supplied to the generator together with the query, and the item is labeled useful if the resulting generation utility exceeds a no-augmentation baseline’s utility ($\delta_j = u_j - u_i > 0$, where u_i is the baseline utility and u_j the

utility with item j added); this makes the labels both task- and generator-dependent, since δ_j is computed with the specific generator being evaluated. Because no official LaMP test split exists, they build the evaluation collection from the first 1,000 entries of the LaMP user-based *development* split, discarding queries with only one useful item (item-side fairness is moot with a single useful candidate). We use their released, generator-specific filtered collections for Flan-T5-Small and Flan-T5-Base directly, with no further subsampling. Because the labels are generator-dependent, the two generators’ retained query sets differ in size (Table 2) and overlap only partially — a given query ID retained for one generator is not guaranteed to be retained for the other.

3.2 Generators

We evaluate two instruction-tuned Flan-T5 generators (Chung et al., 2024; Raffel et al., 2020): Flan-T5-Small (77M parameters) and Flan-T5-Base (250M parameters). Larger Flan-T5 variants exist in the source paper’s evaluation but were not computationally feasible on the hardware available for this project (Section 8). Generation is deterministic beam search (beam size 4, no sampling, no generation-time temperature), with a 128-token generation limit, identical across all seven LaMP tasks including LaMP-3’s numeric-rating task; no task-specific decoding configuration is used.

3.3 Retrievers

We use two retrievers with different underlying representations: a sparse lexical retriever, BM25 (Robertson and Zaragoza, 2009), and an unsupervised dense retriever trained with contrastive learning, Contriever (Izacard et al., 2022) (`facebook/contriever` checkpoint). Both are computed locally by this project’s own pipeline — BM25 via a from-scratch implementation adapted from the `rank_bm25` library, Contriever via the checkpoint above — rather than reused from Kim and Diaz’s released artifacts.

3.4 Rerankers

All three rerankers below operate over the same candidate pool: every item in the query’s profile that the retriever (BM25 or Contriever) scored, which is retrieved and ranked in full before any reranking step. $k = 5$ is the *output truncation size*, not the candidate pool size: each reranked list is truncated to its top 5 items only after the deterministic, MMR, or PL selection procedure has been applied to the full scored profile. For each (generator, retriever) cell, we evaluate 12 reranking settings, each producing a length-5 output list:

- **Deterministic baseline:** the top 5 items of the full scored profile by raw retrieval score, with no fairness or diversity intervention. This is the relevance-only reference condition throughout the paper.
- **MMR**, at seven diversity levels $\lambda \in \{0.15, 0.30, 0.45, 0.55, 0.70, 0.85, 1.00\}$, greedily selecting 5 items from the full scored profile. As a construction check, $\lambda = 1$ should exactly reproduce the deterministic baseline’s outcome (Section 4).
- **PL**, at four fairness-control strengths $\alpha \in \{1, 2, 4, 8\}$ (lower α is a stronger fairness intervention), sampling a length-5 list from the full scored profile via the Plackett–Luce procedure above, with $N = 10$ independently sampled rankings per α per query. Kim and Diaz estimate their stochastic-ranking quantities from $N = 100$ samples; we use $N = 10$ as an explicit, disclosed computational compromise (Section 8).

How the $N = 10$ PL samples become one query-setting row. EU and ILD-Jaccard are computed separately for each of the 10 sampled rankings — one generator call and one diversity computation per sample — and the value reported for a (query, α) setting is the arithmetic mean across those 10. EE-D and EE-R, by contrast, are computed once per query, jointly over the full set of 10 sampled rankings together (the statistically appropriate treatment, since expected exposure is defined over the empirical distribution of rankings a policy produces, not per-sample); no separate per-sample values exist for these two metrics to average. The deterministic baseline and MMR each produce a single ranking, so no aggregation step applies to their settings.

3.5 Metrics

Fairness (EE-D, lower is fairer). We use Kim and Diaz’s own source-normalized EE-D, computed by their vendored implementation of the expected-exposure framework (Diaz et al., 2020) with `normalize=True` (i.e., normalized against theoretical per-query exposure bounds), applied unmodified to our rankings. We do *not* apply any further normalization to EE-D; wherever $\text{EE-D}_{\text{norm}}$ appears in this paper it refers to this same source-normalized quantity, not raw EE-D.

Utility (EU). Task-specific: accuracy for LaMP-1/2, mean absolute error (MAE, lower is better) for LaMP-3, and ROUGE-L for LaMP-4–7.

Diversity (ILD-Jaccard). One minus the mean pairwise Jaccard similarity of the retrieved profile text within a single ranked list; higher is more diverse.

Exposure-relevance alignment (EE-R). Also computed by the vendored expected-exposure implementation with `normalize=True`; used only as a candidate mechanism variable in Section 5.4.

3.6 Normalization

Because LaMP’s seven tasks use three different utility metrics on different scales, raw EU values are not comparable across tasks. We define our own within-query normalized utility, EU_{norm} , distinct from Kim and Diaz’s separate empirical-upper-bound normalization: for query q , generator g , and setting s ,

$$\text{EU}_{\text{norm}}(q, g, s) = \frac{\text{EU}'(q, g, s)}{\max_{s' \in S(q, g)} \text{EU}'(q, g, s')}$$

where

$$\text{EU}'(q, g, s) = \begin{cases} 4.0 - \text{EU}(q, g, s), & \text{LaMP-3} \\ \text{EU}(q, g, s), & \text{otherwise} \end{cases}$$

and $S(q, g)$ is the set of settings evaluated for query q under generator g , *pooled across both retrievers* (up to 24 rows: 2 retrievers $\times$ 12 rerank settings) — our normalization grouping does not condition on retriever, so a query’s denominator is shared by its BM25 and Contriever rows alike, not computed separately per retriever (LaMP-3 uses MAE, where lower is better, so it is flipped before normalizing; every other task’s raw EU is used as-is). This is true max-normalization (division by the largest raw value a query achieved under that generator, across both retrievers), not min-max scaling. ILD_{norm} and $\text{EE-R}_{\text{norm}}$ are defined identically, dividing by each (q, g) pair’s own maximum raw ILD or EE-R value across $S(q, g)$. $\text{EE-D}_{\text{norm}}$ is the one exception: it is never further normalized, since it is already the source-normalized quantity described above.

The normalization denominator is exactly zero, and EU_{norm} undefined, precisely when a query scores raw $\text{EU} = 0$ under *every* setting in $S(q, g)$ — verified directly for all 237 affected queries, with no exceptions. This excludes 5,688 of 163,392 pre-exclusion query-setting rows (3.5%), concentrated in LaMP-1, 2, 5, and 6 (zero excluded rows in LaMP-3 and LaMP-7); these are queries on which no

method received any credit, not queries where methods tied at an informative nonzero value. As a sensitivity check, we re-ran the headline specification for each of Findings 2–4 with these rows filled at $EU_{\text{norm}} = 0$ rather than excluded: coefficients moved modestly and no conclusion changed sign or significance status (Finding 2, pooled coef. EE-D: $+0.0735 \rightarrow +0.0697$, $p < 10^{-50}$; Finding 3, pooled coef. ILD: $-0.0675 \rightarrow -0.0838$, $p = 0.012 \rightarrow 0.0017$; Finding 4, AME from the interaction model — the headline quantity of Section 5.4, not the additive coefficient: $-0.0192 \rightarrow -0.0188$, $p = 1.5 \times 10^{-33} \rightarrow 3.1 \times 10^{-34}$). The primary analysis throughout this paper excludes these rows because EU_{norm} is undefined for them; the sensitivity check above confirms this exclusion is not driving any headline conclusion.

3.7 Statistical Methodology

All hypothesis tests in this paper are two-sided, at a nominal significance level of 0.05 (written out as “significance level” rather than α throughout this section, to avoid clashing with PL’s fairness-control parameter α). The clustering and grouping unit used throughout, unless stated otherwise, is the query identifier (`lamp_num+qid`).

Paired tests. Wilcoxon signed-rank tests on per-query EU_{norm} , paired against the deterministic baseline, Holm-corrected within each (generator, retriever) cell’s test family.

Cluster-robust inference. Because each query is observed under multiple settings (every α , every λ , both methods), we report CR1 sandwich-estimator standard errors clustered by query wherever a pooled regression is the primary estimate, alongside the naive (non-clustered) p -value for comparison where informative.

Leave-one-task-out (LOTO). For every headline pooled coefficient, we re-fit the identical specification after dropping each of the seven LaMP tasks in turn, checking whether the coefficient’s sign or significance status changes.

Fixed effects. To separate within-query treatment variation (e.g., α changing for the same query) from between-query composition (some queries are systematically easier, fairer-scoring, or more diverse regardless of setting), we estimate two nested fixed-effects specifications via the within (demeaning) estimator with cluster-robust standard errors: a *query* fixed effect (one dummy per task $\times$ query id), and a finer *query* $\times$ *generator* $\times$ *retriever* fixed effect. Query IDs are shared across generators but the retained sets differ (Section 3), so for query IDs retained under more than one generator/retriever cell, the coarser per-query fixed effect can still let that cross-cell composition contribute to the “within” variation it attributes to the treatment dial; the finer fixed effect restricts identification to variation across α/λ /method settings within each exact observed query-generator-retriever cell, avoiding reliance on cross-generator or cross-retriever composition, and does not require every query to appear under every cell.

Interaction and marginal-effects model. For Finding 4 (Section 5.4), we test whether a single additive PL-vs-MMR coefficient is an adequate summary by adding a $PL \times ILD_{\text{norm}}$ interaction term. When the interaction is non-negligible, we report the conditional PL–MMR gap at several points across the observed ILD_{norm} range, and the average marginal effect (AME) — the interaction model evaluated at the empirical mean of ILD_{norm} — with delta-method standard errors computed from the full coefficient covariance matrix.

3.8 Reproducibility

A single fixed run-level random seed (default 42) is applied identically before every query’s Plackett–Luce/Gumbel-max sampling, and also seeds the language model, PyTorch, and CUDA random number generators (with deterministic cuDNN operations enabled). This is a reproducible, not a

per-query-derived, seed: the same seed value is re-applied before each query’s sampling step, rather than each query drawing from a uniquely-derived seed. PL’s rankings remain genuinely stochastic — a fixed seed makes that randomness reproducible across re-runs, not absent.

3.9 Task Weighting

LaMP task sample sizes differ substantially after filtering (Table 2). Pooled query-level regressions and tests throughout this paper weight each *query* equally, not each task; larger tasks therefore contribute proportionally more observations to any pooled estimate. We report task-level heterogeneity and leave-one-task-out results alongside every pooled number specifically so that this implicit weighting does not go unnoticed.

4 Experimental Setup

Table 1 summarizes the full factorial design: 2 generators $\times$ 2 retrievers $\times$ 12 reranking settings $\times$ 7 LaMP tasks, for 336 total experiment configurations and 163,392 pre-exclusion query-setting rows before the normalization exclusion of Section 3.6. Table 2 gives the retained query count per task and generator after the source paper’s per-generator utility-label filtering.

Table 1: Experimental design.

Component	Values
Benchmark	LaMP, all 7 tasks, “user” split
Generators	Flan-T5-Small (77M), Flan-T5-Base (250M)
Retrievers	BM25, Contriever
Candidate pool	full retriever-scored profile (per query)
Output/truncation size	$k = 5$
Deterministic baseline	top- k of the full scored profile
MMR	$\lambda \in \{0.15, 0.30, 0.45, 0.55, 0.70, 0.85, 1.00\}$
PL	$\alpha \in \{1, 2, 4, 8\}$, $N = 10$ sampled rankings per α
Settings per cell	12 (1 deterministic + 7 MMR + 4 PL)
Total configurations	336
Pre-exclusion query-setting rows	163,392

Table 2: Retained query count per LaMP task and generator, after the source paper’s per-generator utility-label filtering.

LaMP task	Flan-T5-Base	Flan-T5-Small
1	232	51
2	280	192
3	378	311
4	827	833
5	759	826
6	783	760
7	211	365

Sanity check: MMR at $\lambda = 1$. By construction, $\lambda = 1$ removes MMR’s diversity term entirely and should reproduce the deterministic baseline. We verify this directly rather than assuming it: the maximum paired EU difference between the two conditions is **exactly 0.0 across all 13,616 paired queries**. This confirms outcome equality, not identical document rankings, which we did not independently verify at the document-ID level.

4.1 Premise Check: Does PL’s Fairness Mechanism Actually Increase Diversity Here?

The diversity-as-mechanism hypothesis requires that PL’s fairness intervention actually moves diversity as a byproduct; we test this rather than assume it. Figure 1 shows both quantities as PL’s α varies, averaged across the four (generator, retriever) cells.

Query-level coupling is weak. Within PL rows, the Pearson correlation between $\text{EE-D}_{\text{norm}}$ and ILD_{norm} is $r = -0.086$ ($n = 54,464$) — a small effect once individual queries are considered, amid substantial per-query variation in both quantities.

The aggregate, by- α pattern is monotonic. Table 3 reports the mean source-normalized EE-D and mean raw ILD-Jaccard at each α : as α increases (fairness intervention weakens toward the deterministic baseline), disparity rises and diversity falls, monotonically in the expected directions.

Table 3: Aggregate fairness and diversity by PL α (lower α is fairer).

PL α	Mean source-normalized EE-D	Mean ILD-Jaccard (raw)
1	0.210	0.879
2	0.248	0.870
4	0.657	0.838
8	0.948	0.830

The premise therefore holds *in aggregate, as a tendency across the fairness dial*, but not as a tight per-query coupling. We report both because this looseness matters downstream: it is part of why a PL-vs-MMR comparison adjusted for measured diversity (Section 5.4) cannot be treated as a clean causal isolation of fairness’s own contribution.

In our experiments, stronger PL fairness intervention coincides with higher diversity in aggregate; MMR changes diversity with no fairness objective




Figure 1: In our experiments, stronger PL fairness intervention coincides with higher diversity as an aggregate, by- α tendency (left panel): PL’s dial moves both expected exposure disparity and diversity at once. This aggregate coupling is not tight at the level of individual queries (row-level Pearson $r = -0.086$, Section 4.1). MMR’s dial (right panel) changes measured diversity without an explicit fairness objective, with exposure disparity remaining near the deterministic level (EE-D ≈ 1), which is what makes it usable as a diversity-focused comparator against PL. Bold lines are the mean across the four (generator, retriever) cells; thin lines are individual cells.

5 Results

5.1 Finding 1: PL Reranking Generally Does Not Improve Utility

We test each PL α against the deterministic baseline, paired per-query, Holm-corrected within each cell’s 4-test family (4α values $\times$ 4 generator/retriever cells = 16 tests total). This scopes the test specifically to PL, since MMR has no fairness objective and does not belong in a test about whether fairness pays off; MMR-vs-baseline comparisons are evidence about diversity instead (Section 5.3).

Of the 16 tests, **10 are statistically significant, and 8 of those are losses**. The two significant gains are both the same PL configuration ($\alpha = 4$) in two different cells, and are small ($+0.0014$ and $+0.0025$ in EU_{norm}) against losses roughly an order of magnitude larger elsewhere.

Kim and Diaz’s own reported result is more nuanced than “fair rankings often outperform”: they report an overall fairness/quality tradeoff in which moderate fairness settings sometimes preserve or improve utility. Given that our evaluation covers 2 of their 4 generators, 2 of their 3 retrievers, and $N = 10$ sampled rankings against their $N = 100$, we do not frame this as contradicting their result outright. Within the two smaller generators and two retrievers evaluated here, however, we do not reproduce the pattern that moderate fair reranking frequently preserves or improves deterministic-baseline utility: the direction is overwhelmingly negative, with two small, real, but practically negligible exceptions. Figure 2 shows the deterministic baseline’s normalized utility across tasks and cells, illustrating the substantial cross-task differences under our within-query normalization that motivate reporting EU_{norm} rather than raw EU throughout.




Figure 2: Deterministic-baseline normalized utility (EU_{norm}) by LaMP task, for all four (generator, retriever) cells. EU_{norm} is each query’s raw utility divided by the largest raw utility it achieved across its settings under that generator (Section 3.6), so this figure shows the deterministic condition’s relative performance under that within-query normalization, not an absolute or task-independent utility scale — it illustrates why per-query normalization is necessary for any cross-task comparison in this paper, not what generation quality is achievable in some absolute sense.

5.2 Finding 2: A Real Fairness–Utility Tradeoff Within PL’s Own Fairness Dial

We regress EU_{norm} on $EE\text{-}D_{\text{norm}}$, restricted to PL rows only, controlling for ILD_{norm} and generator/retriever indicators:

$$\begin{aligned} EU_{\text{norm}}(q, g, r, s) = & \beta_0 + \beta_1 EE\text{-}D_{\text{norm}}(q, g, r, s) \\ & + \beta_2 ILD_{\text{norm}}(q, g, r, s) + \beta_3 \mathbf{1}\{\text{Base}\} + \beta_4 \mathbf{1}\{\text{Contriever}\} \\ & + \varepsilon(q, g, r, s), \quad s \in \text{PL rows.} \end{aligned} \quad (1)$$

Restricting to PL rows avoids mixing a within-method fairness manipulation (PL’s α) with a between-method contrast (MMR and the deterministic baseline sit at a fixed $EE\text{-}D \approx 1$ regardless of λ), which would not be a clean fairness-dial estimate.

The primary estimate is $\hat{\beta}_1 = +0.0735$ (cluster-robust $p = 6.3 \times 10^{-54}$, $n = 52,568$, 4,269 query clusters; the naive, non-clustered $p = 2.1 \times 10^{-93}$ overstates confidence by treating repeated observations of the same query as independent). Because lower $EE\text{-}D$ means fairer, a *positive* coefficient means that **less fair, more disparate rankings are associated with higher utility** — the opposite of what a “fairness helps” story would predict.

Robustness. Leave-one-task-out, on the identical specification and cluster-robust inference, keeps the coefficient positive and highly significant under every single-task exclusion, ranging from $+0.0613$ (dropping LaMP-4, $p = 1.9 \times 10^{-27}$) to $+0.0834$ (dropping LaMP-5, $p = 2.7 \times 10^{-53}$); no sign or significance flips. Query fixed effects — isolating the within-query association between $EE\text{-}D_{\text{norm}}$ and EU_{norm} from between-query composition — give $+0.0518$ ($p = 2.0 \times 10^{-30}$); the finer query $\times$ generator $\times$ retriever fixed effect, which isolates purely the α -driven variation for one exact query in one exact cell, gives $+0.0526$ ($p = 2.4 \times 10^{-30}$), essentially unchanged from the coarser fixed effect. The pooled estimate was picking up some between-query variation, but the within-query association remains substantial under every specification we tested (Table 8).

Task heterogeneity. Fit on each task alone (no leave-one-task-out exclusion, no other controls), the coefficient is positive in all seven tasks, but not statistically significant on two of them in isolation (LaMP-3, $p = 0.31$; LaMP-7, $p = 0.71$); the remaining five are individually significant and positive. Unlike Findings 3 and 4, no task individually reverses this association’s sign. We nonetheless describe this as a robust *pooled* association rather than a claim that it is statistically detectable within every single task in isolation, since two of the seven do not reach significance on their own.

5.3 Finding 3: No General Utility Benefit From MMR Diversity

If measured $ILD\text{-}Jaccard$ were sufficient to explain PL’s apparent utility behavior, variation induced by MMR in that measured quantity should provide corresponding evidence of a general utility benefit. MMR has no exposure-fairness objective, but changing λ also changes the relevance–diversity tradeoff, so this comparison is informative without being a causal isolation of diversity. We regress EU_{norm} on ILD_{norm} within MMR rows only (exposure disparity is fixed near the deterministic level, $EE\text{-}D \approx 1$, for all MMR settings):

$$\begin{aligned} EU_{\text{norm}}(q, g, r, s) = & \gamma_0 + \gamma_1 ILD_{\text{norm}}(q, g, r, s) + \gamma_2 \mathbf{1}\{\text{Base}\} + \gamma_3 \mathbf{1}\{\text{Contriever}\} + \varepsilon(q, g, r, s), \\ & s \in \text{MMR rows.} \end{aligned} \quad (2)$$

We find **no evidence of a general utility benefit from increased MMR diversity**: the pooled association is modestly negative, $\hat{\gamma}_1 = -0.0675$ ($n = 91,994$). The cluster-robust $p = 0.0115$ provides modest evidence for that association, rather than overwhelming evidence — a considerably

Overview: fairness vs. utility, all methods (per-setting means) - PL's α -trajectory highlighted; the statistical test (Finding 2) uses PL rows only




Figure 3: Fairness ($EE-D_{\text{norm}}$, higher is less fair) versus utility (EU_{norm}), per-setting means, one panel per generator/retriever cell. PL's α -trajectory is highlighted (triangles); MMR (circles) and the deterministic baseline (star) are shown for context, but the Finding 2 regression (Eq. 1) uses PL rows only.

more modest significance level than the naive $p = 5.6 \times 10^{-10}$ suggested, since the repeated-query structure was inflating confidence.

Robustness. The pooled coefficient stays negative under every single-task exclusion (a sign check; no leave-one-task-out significance claim is at stake here), but it does not hold in every task individually: fit on LaMP-2 or LaMP-5 alone, the coefficient is positive. The pooled negative association is robust to removing any one task, while the direction itself is heterogeneous across tasks — the diversity–utility association is benchmark-dependent/heterogeneous across tasks, not a fixed property of RAG generation. Query fixed effects give -0.0820 ($p = 2.4 \times 10^{-4}$), if anything larger in magnitude than the pooled estimate; the finer query $\times$ generator $\times$ retriever fixed effect gives -0.0785 ($p = 5.0 \times 10^{-4}$), essentially unchanged (Table 8).

Per-cell heterogeneity. Figure 4 and Table 4 report the identical bivariate specification (ILD_{norm} only, no controls) fit separately within each cell, with cluster-robust inference. Only the Flan-T5-Base/BM25 cell is individually significant. A formal test of whether the visually larger Base-cell coefficients reflect a genuine generator-size interaction — a pooled within-MMR regression with an ILD_{norm} $\times$ generator interaction term, cluster-robust by query and controlling for task and retriever — gives $p = 0.28$: not significant. The descriptive per-cell pattern in Figure 4 is real and visible, but the formal test does not confirm generator size as its driver.

Table 4: Per-cell bivariate regression of EU_{norm} on ILD_{norm} within MMR rows (no controls), matching Figure 4 exactly.

Cell	Coefficient	p
Flan-T5-Small / BM25	-0.0374	0.34
Flan-T5-Small / Contriever	-0.0149	0.69
Flan-T5-Base / BM25	-0.1483	8.9×10^{-5}
Flan-T5-Base / Contriever	-0.0728	0.068




Figure 4: Within-MMR association between diversity and utility, per cell, with exposure disparity fixed at the deterministic level. Coefficients correspond to Table 4.

Direct MMR-vs-deterministic comparison. As a complementary, non-regression check, we test whether MMR ever significantly beats or loses to the deterministic baseline directly: paired per-query, all 7 λ values $\times$ 4 cells = 28 comparisons (every MMR setting run, not a pre-selected subset), Holm-corrected within each cell’s 7-test family. **4 of 28 are significant, and all 4 are losses** (Table 5). The pattern is consistent with the negative diversity–utility association above (lower λ trades more relevance for more diversity, not diversity in isolation, so this is not independent confirmation of a diversity-specific effect): every significant loss occurs at the two strongest-diversity settings, and the per-cell delta shrinks toward zero monotonically as $\lambda \rightarrow 1$, where it is exactly 0.0 in every cell by the construction check of Section 4. Overall, 19 of 28 comparisons point negative, and MMR shows no significant utility *win* over the deterministic baseline in any of the 28 settings tested.

Table 5: MMR versus deterministic baseline, paired per-query, aggregated over the four generator/retriever cells (28 comparisons total, Holm-corrected within each cell’s 7-test family).

MMR λ	Cells with negative Δ (of 4)	Cells significant (of 4)
0.15	4	3
0.30	4	1
0.45	4	0
0.55	2	0
0.70	3	0
0.85	2	0
1.00	0 (exactly zero, by construction)	0

5.4 Finding 4: A Residual PL–MMR Utility Difference After Adjusting for Diversity

Restricting to PL and MMR rows, we ask whether being a PL setting (versus MMR) predicts a further utility difference once measured diversity is accounted for. This is a regression adjustment, not a matching procedure.

Additive specification. Table 6 reports the per-cell coefficient of a PL indicator on EU_{norm} , controlling for ILD_{norm} ; all four cells are negative, significant, and of similar magnitude. Pooling with a single common ILD_{norm} slope for PL and MMR gives an additive PL-indicator coefficient of -0.0194 (cluster-robust $p = 8.8 \times 10^{-34}$) — but before treating this single number as the answer, we test whether a single common slope is the right model.

Table 6: Per-cell coefficient of the PL indicator on EU_{norm} , controlling for ILD_{norm} (additive specification).

Cell	Coefficient	p (cluster-robust)
Flan-T5-Small / BM25	-0.0130	1.1×10^{-7}
Flan-T5-Base / BM25	-0.0226	5.5×10^{-17}
Flan-T5-Small / Contriever	-0.0186	1.4×10^{-12}
Flan-T5-Base / Contriever	-0.0227	1.9×10^{-16}

Primary specification: does the gap depend on diversity level? We add a $PL \times ILD_{\text{norm}}$

interaction term:

$$\begin{aligned} \text{EU}_{\text{norm}}(q, g, r, s) = & \beta_0 + \beta_1 \mathbf{1}\{\text{PL}\}(s) + \beta_2 \text{ILD}_{\text{norm}}(q, g, r, s) \\ & + \beta_3 (\mathbf{1}\{\text{PL}\}(s) \times \text{ILD}_{\text{norm}}(q, g, r, s)) \\ & + \beta_4 \mathbf{1}\{\text{Base}\} + \beta_5 \mathbf{1}\{\text{Contriever}\} + \sum_t \beta_t \mathbf{1}\{\text{task} = t\} + \varepsilon(q, g, r, s). \end{aligned} \quad (3)$$

Controls are a generator indicator, a retriever indicator, and LaMP-task fixed effects. Inference is cluster-robust by query using the CR1 sandwich estimator; Table 6’s per-cell p -values use the same cluster-robust inference rather than classical standard errors. The interaction is significant in the pooled specification ($\hat{\beta}_3 = +0.0293$, $p = 0.0125$): the two methods’ ILD_{norm} slopes are not identical, so a single additive coefficient understates what is happening. Because ILD_{norm} is not centered, $\hat{\beta}_1$ alone is the model’s fitted PL–MMR gap specifically at $\text{ILD}_{\text{norm}} = 0$; that value does occur in the data (both methods’ observed ILD_{norm} minimum is exactly 0.0) but is rare, since ILD_{norm} ’s own 10th percentile is already 0.84. Table 7 reports the conditional gap, $\Delta(\text{ILD}_{\text{norm}}) = \hat{\beta}_1 + \hat{\beta}_3 \cdot \text{ILD}_{\text{norm}}$, at several points across the observed range, together with the average marginal effect (AME), the same quantity evaluated at the empirical mean of ILD_{norm} .

Table 7: Conditional PL–MMR gap at points across the observed ILD_{norm} range, and the average marginal effect (AME), from the interaction model (Eq. 3).

ILD_{norm} level	$\Delta(\text{PL} - \text{MMR})$	95% CI	p
0 (observed floor, rare)	−0.0464	[−0.068, −0.025]	2.1×10^{-5}
p_{10} (0.84)	−0.0218	[−0.025, −0.018]	2.5×10^{-32}
median (0.96)	−0.0181	[−0.021, −0.015]	2.8×10^{-27}
p_{90} (0.99)	−0.0172	[−0.021, −0.014]	3.6×10^{-21}
mean (AME)	−0.0192	[−0.022, −0.016]	1.5×10^{-33}

The gap is largest, about −0.046, in the small, rare tail of low-diversity settings, and narrows to roughly −0.017 to −0.022 across the bulk of the data (p_{10} to p_{90} of observed ILD_{norm}). We treat the AME, −0.0192, as the headline single-number summary of Finding 4: a proper average over the interaction model rather than an assumption of one fixed slope, and it lands almost exactly where the simpler additive coefficient (−0.0194) did.

Is the interaction itself robust? Not entirely, though the AME built from it is. Under the strictest identification we use — query×generator×retriever fixed effects — the interaction shrinks and loses significance ($\hat{\beta}_3 = +0.0157$, $p = 0.15$), about half the pooled magnitude and no longer distinguishable from zero. Leave-one-task-out on the pooled interaction model keeps the coefficient positive when 6 of 7 tasks are dropped individually (range +0.023 to +0.037; five of those six drops are individually significant, p from 0.003 to 0.021, and the sixth, dropping LaMP-7, is not, $p = 0.061$), but **flips sign when LaMP-6 is dropped** (−0.0280, $p = 0.16$, itself not significant either way) — the only one of the seven folds to do so. We therefore treat the diversity-dependence pattern — the interaction term itself — as a suggestive, secondary observation, not an established one. What *does* survive, robustly, across every specification we tried (pooled, query fixed effects, query×generator×retriever fixed effects, and all seven leave-one-task-out folds) is the AME itself: it ranges narrowly from −0.0161 to −0.0217 across every leave-one-task-out fold (all $p < 10^{-17}$), and from −0.0192 to −0.0200 across the pooled and fixed-effects variants (Table 8), always highly significant and never close to flipping sign. Figure 5 shows the raw per-quintile pattern that motivates the interaction check; we caution against reading the apparent low-vs-high-diversity gap

in that figure as an established, robust pattern in its own right — it is the AME in Table 8, not the bin-level shape of Figure 5, that is the robust quantity.

What this does and does not establish. Adjusting for ILD_{norm} controls for one measured outcome of the ranking, not for everything that differs between PL and MMR: PL is stochastic and MMR deterministic, the two methods may preserve the underlying relevance ordering differently, and PL introduces variance across repeated draws that MMR does not have. The AME is therefore an *average adjusted difference* between PL and MMR after controlling for measured ILD_{norm} , not a causal estimate of what fairness itself costs.

Common support. Does the comparison rely on extrapolating outside the diversity range either method actually reaches? No: PL’s and MMR’s measured ILD_{norm} distributions essentially fully overlap (closely aligned 5th/median/95th percentiles), and restricting to their common support changes nothing: the secondary additive PL-indicator coefficient is -0.0193 ($p = 8.8 \times 10^{-34}$; all 144,562 rows already fall inside the overlap) — the interaction-model AME above is grounded in the full sample, unaffected by this restriction.

Mechanism test. We hypothesized that PL pays a *relevance* cost — stochastic sampling moves away from the highest-ranked item, and that specifically, not diversity, explains the residual. EE-R is a direct proxy for exposure-relevance alignment. In this secondary additive mechanism-check specification, adding $EE-R_{\text{norm}}$ as a covariate moves the PL-indicator coefficient from -0.0194 to -0.0187 ($p = 3.0 \times 10^{-34}$) — essentially unchanged, even though $EE-R_{\text{norm}}$ itself is strongly associated with utility on its own ($\hat{\beta} = +0.158$, $p = 5.1 \times 10^{-104}$, a sanity check that $EE-R_{\text{norm}}$ behaves as expected). **This relevance-cost hypothesis is not confirmed by the test.** PL underperforms MMR on average after adjusting for ILD_{norm} , robustly across the specifications tested; the pooled interaction model suggests a larger disadvantage in the rare low-diversity tail, but that interaction does not survive the strictest fixed-effects and leave-one-task-out checks, and the residual mechanism remains unidentified. It may relate to the stochasticity itself (variance, not just a mean shift) or to some other ranking property that ILD_{norm} and $EE-R_{\text{norm}}$ together do not capture, but the present experiment does not identify which. Supplementary Figure 7 (Appendix A) provides the corresponding per-task descriptive PL–MMR differences within shared ILD bins; we present it as descriptive detail underlying the AME above, not as independent robustness evidence.

Utility by diversity level: MMR vs PL vs the original deterministic list
(task-aligned bins: only (task, bin) cells where both methods have data)




Figure 5: Mean utility by diversity quintile, MMR versus PL, with the deterministic baseline shown for reference. Error bars are 95% confidence intervals from a task-stratified bootstrap (400 resamples, each task resampled with replacement and averaged with equal task weight, 2.5th/97.5th percentiles). This raw, descriptive pattern motivated the interaction check of Table 7; as discussed in the text, the AME is the robust summary of Finding 4, not the specific low-vs-high quintile shape visible here.

6 Robustness Analyses

Leave-one-task-out. The headline pooled coefficients for Findings 2 and 3 were checked leave-one-task-out; neither’s main conclusion depends on any single task remaining in the sample. Finding 4’s AME is similarly leave-one-task-out robust (every fold -0.016 to -0.022 , always significant), while the $\text{PL} \times \text{ILD}_{\text{norm}}$ interaction underlying its secondary diversity-dependence framing is not (it flips sign when LaMP-6 is dropped). This does not cover every secondary specification in this paper (per-cell regressions, the generator interaction, the EE-R mechanism check); those were not individually re-run leave-one-task-out.

Cluster-robust inference. Added throughout to correct for repeated observations of the same query across settings; it changes exact significance levels (most notably Finding 3, where the naive $p = 5.6 \times 10^{-10}$ becomes $p = 0.0115$ once clustering is applied) but does not overturn any finding’s direction or overall conclusion.

Fixed effects. Table 8 summarizes all three findings under query fixed effects and the finer query $\times$ generator $\times$ retriever fixed effect. For Finding 4, we report the AME from the interaction model at each fixed-effects level (not the additive coefficient) as the headline robustness quantity, per the primary specification of Section 5.4; the additive coefficient (also fixed-effects-robust, at -0.0200 under both levels) is a secondary check. The two fixed-effects levels barely differ from each other for any finding, indicating that cell composition was not doing hidden work in the coarser, query-only version.

Table 8: Pooled coefficients versus query fixed effects and the finer query $\times$ generator $\times$ retriever fixed effect. Finding 4’s row reports the AME from the interaction model (Eq. 3), not the additive coefficient, consistent with Section 5.4.

Finding	Pooled	Query FE	Query $\times$ gen. $\times$ retr. FE
2 (coef. EE-D, PL rows)	+0.0735 ($p = 6.3 \times 10^{-54}$)	+0.0518 ($p = 2.0 \times 10^{-30}$)	+0.0526 ($p = 2.4 \times 10^{-30}$)
3 (coef. ILD, MMR rows)	-0.0675 ($p = 0.0115$)	-0.0820 ($p = 2.4 \times 10^{-4}$)	-0.0785 ($p = 5.0 \times 10^{-4}$)
4 (AME, PL+MMR)	-0.0192 ($p = 1.5 \times 10^{-33}$)	-0.0200 ($p = 1.6 \times 10^{-36}$)	-0.0200 ($p = 2.0 \times 10^{-36}$)

Normalization sensitivity. As described in Section 3.6, filling the 3.5% of rows excluded by the zero-denominator condition, rather than dropping them, leaves all three conclusions unchanged.

Task-robustness figure. Figure 6 visualizes the leave-one-task-out check for Findings 2 and 3’s pooled coefficients together, each on its exact final specification (Finding 2: PL rows only; Finding 3: MMR rows only); both retain their sign and significance status under every single-task exclusion. Finding 4’s leave-one-task-out robustness is reported separately in Section 5.4 and Table 8, since its headline quantity (the interaction-model AME) is not a single additive coefficient this panel layout can show alongside Findings 2–3 without visually implying the fragile $\text{PL} \times \text{ILD}_{\text{norm}}$ interaction is itself the headline result.

Task-robustness, Findings 2-3: both pooled coefficients keep their sign and significance under every single-task exclusion (red = significance lost when dropped)




Figure 6: Leave-one-task-out robustness for Findings 2 and 3's pooled coefficients. Top row: coefficient when each task is dropped, against the full-sample estimate (blue band, 95% CI). Bottom row: coefficient fit on each task alone — Finding 2 (left) is positive in every task alone, though not statistically significant in two of them; Finding 3 (right) reverses sign in two tasks fit alone. No task exclusion flips the sign or significance status of either pooled coefficient.

7 Discussion

This paper set out to answer a specific question: when fair reranking appears to help or hurt RAG generation quality, is that attributable to fairness itself, or to the within-list diversity the fairness mechanism happens to induce? Across the settings we test, the answer is that **measured within-list diversity (ILD-Jaccard) alone does not account for the difference between PL and MMR utility**. PL’s own fairness dial is associated with a real, robust, and negative relationship to utility (Finding 2); MMR’s diversity dial shows no general utility benefit of its own, and if anything a modest negative association (Finding 3); and PL retains an average utility disadvantage relative to MMR even after adjusting for measured diversity (Finding 4). If measured ILD-Jaccard were sufficient to explain PL’s apparent utility behavior, we would expect the diversity variation MMR induces in that measured quantity to show a corresponding general utility benefit on its own; we do not observe such a pattern.

This is a narrower conclusion than Kim and Diaz’s original finding might suggest, but the two results are not necessarily in tension: our evaluation covers 2 of their 4 generators, 2 of their 3 retrievers, and a tenfold-lower PL sample count ($N = 10$ versus $N = 100$), and we do not claim to have reproduced their full experimental grid. What we can say is that, within the Flan-T5-Small and Flan-T5-Base, BM25 and Contriever setting we do test, fair reranking does not reproduce the pattern of frequently preserving or improving utility, and that the diversity-as-mechanism hypothesis specifically does not explain the residual gap between the two reranking families.

It is important to be precise about what the PL-vs-MMR comparison in Finding 4 does and does not identify. Because PL and MMR differ along several dimensions besides measured diversity — stochastic sampling versus deterministic selection, and how each method preserves or disturbs the underlying relevance ordering, among others — the residual utility gap after adjusting for ILD_{norm} is an *adjusted association*, not a causal estimate of what fairness itself costs. We tested the most obvious candidate explanation directly: that PL’s randomization imposes a relevance cost, measurable via exposure-relevance alignment (EE-R). Adding EE-R as a covariate left the residual essentially unchanged, so this specific mechanism is not confirmed. The residual mechanism remains unidentified; it may relate to the variance PL’s stochastic sampling introduces, rather than only its mean shift, or to some property of the ranking that neither ILD_{norm} nor $EE-R_{\text{norm}}$ captures, but this remains an open question for future work, which we discuss should require a design that varies fairness and diversity independently rather than only adjusting for one while measuring the other.

Finally, task heterogeneity recurs throughout, though not always as a sign reversal: Finding 2’s association is positive in all seven tasks fit alone but not statistically distinguishable from zero in two of them (LaMP-3, LaMP-7); Finding 3’s association reverses sign on LaMP-2 and LaMP-5 alone; Finding 4’s secondary diversity-dependence interaction reverses sign when LaMP-6 is excluded. None of these task-level patterns overturn the corresponding pooled, leave-one-task-out-robust conclusion, but they are a reminder that these associations are properties of a query-weighted pool of seven quite different personalization tasks, not of RAG generation in general.

8 Limitations

- **Generator scale.** We evaluate only Flan-T5-Small (77M) and Flan-T5-Base (250M). Whether this tradeoff grows, shrinks, or reverses at larger generator scales is untested here; utility labels exist for a much larger (XXL) generator in the source release, but it does not fit the GPU hardware available for this project.
- **Retriever coverage.** We evaluate BM25 and Contriever only, 2 of the 3 retrievers in the

source paper’s evaluation.

- **PL sample count.** $N = 10$ sampled rankings per (q, α) pair, against the source paper’s $N = 100$; our PL-based EE-D/EU estimates carry more Monte Carlo noise than theirs. A cheap follow-up — recomputing EE-D only at $N = 100$ on a subset, without regenerating text — would quantify Monte Carlo sensitivity of the exposure metrics specifically, but would not by itself bound Monte Carlo uncertainty in generated utility; this was not run for this paper.
- **Diversity operationalization.** We measure diversity as ILD-Jaccard, a single specific notion of within-list textual variety; our conclusion that measured diversity does not account for the PL–MMR gap is a claim about *this* operationalization, not about every possible notion of diversity.
- **Loose fairness–diversity coupling.** EE-D and ILD are conceptually distinct quantities that are only loosely, not strictly, coupled under PL in this data (Section 4.1); this limits how cleanly any ILD_{norm} -adjusted comparison can isolate fairness’s own contribution.
- **PL and MMR differ beyond diversity.** PL is stochastic and MMR deterministic; the two methods may preserve the original relevance ordering differently, and only PL introduces sampling variance across repeated draws. Finding 4’s residual should not be read as an isolated measurement of fairness’s causal effect for this reason.
- **Task dependence.** Findings 2 and 3’s pooled effects are leave-one-task-out robust. Finding 2’s association is positive in every task fit alone, though not statistically significant in two of them (LaMP-3, LaMP-7); Finding 3’s association does reverse sign on two tasks fit alone (LaMP-2, LaMP-5). Finding 4’s AME is similarly robust, but the interaction term underlying its secondary diversity-dependence framing is not (it loses significance under the strictest fixed effects and flips sign when LaMP-6 is dropped); we treat that pattern as secondary for exactly this reason.

9 Conclusion

We asked whether an apparent fairness benefit in RAG generation quality is attributable to fairness itself or to the within-list diversity that a fairness mechanism happens to induce. Using all available queries retained by Kim and Diaz’s generator-specific fairness-evaluation collections across all seven LaMP tasks, two generators, and two retrievers, we find that PL fair reranking does not generally improve utility over a deterministic baseline; that within PL’s own fairness dial, greater exposure equality is robustly associated with lower utility; that MMR’s explicit diversity dial shows no general utility benefit on its own; and that PL retains an average utility disadvantage relative to MMR even after adjusting for measured diversity, a disadvantage that a direct test of a relevance-cost mechanism does not explain. Taken together, **measured within-list diversity (ILD-Jaccard) alone does not account for the PL–MMR utility difference** we observe, while the residual itself should not be read as the isolated causal cost of fairness, since PL and MMR differ along dimensions besides diversity. Determining what does explain that residual — and whether these patterns persist at larger generator scale and across additional retrievers — is the natural next step, and would benefit from an experimental design that varies fairness and diversity independently rather than adjusting for one while measuring the other.

References

- Asia J. Biega, Krishna P. Gummadi, and Gerhard Weikum. Equity of attention: Amortizing individual fairness in rankings. In *The 41st International ACM SIGIR Conference on Research & Development in Information Retrieval (SIGIR '18)*, pages 405–414. ACM, 2018. doi: 10.1145/3209978.3210063.
- Jaime Carbonell and Jade Goldstein. The use of MMR, diversity-based reranking for reordering documents and producing summaries. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '98)*, pages 335–336, New York, NY, USA, 1998. ACM.
- Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Alex Castro-Ros, Marie Pellat, Kevin Robinson, Dasha Valter, Sharan Narang, Gaurav Mishra, Adams Yu, Vincent Zhao, Yanping Huang, Andrew Dai, Hongkun Yu, Slav Petrov, Ed H. Chi, Jeff Dean, Jacob Devlin, Adam Roberts, Denny Zhou, Quoc V. Le, and Jason Wei. Scaling instruction-finetuned language models. *Journal of Machine Learning Research*, 25(70):1–53, 2024. arXiv:2210.11416.
- Fernando Diaz, Bhaskar Mitra, Michael D. Ekstrand, Asia J. Biega, and Ben Carterette. Evaluating stochastic rankings with expected exposure. In *Proceedings of the 29th ACM International Conference on Information and Knowledge Management (CIKM '20)*, pages 275–284. ACM, 2020. doi: 10.1145/3340531.3411962.
- Gautier Izacard, Mathilde Caron, Lucas Hosseini, Sebastian Riedel, Piotr Bojanowski, Armand Joulin, and Edouard Grave. Unsupervised dense information retrieval with contrastive learning. *Transactions on Machine Learning Research*, 2022. arXiv:2112.09118.
- To Eun Kim and Fernando Diaz. Towards fair RAG: On the impact of fair ranking in retrieval-augmented generation. In *Proceedings of the 2025 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval (ICTIR '25)*. ACM, 2025. doi: 10.1145/3731120.3744599. arXiv:2409.11598.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*, 2020. arXiv:2005.11401.
- R. Duncan Luce. *Individual Choice Behavior: A Theoretical Analysis*. John Wiley & Sons, New York, 1959.
- R. L. Plackett. The analysis of permutations. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 24(2):193–202, 1975. doi: 10.2307/2346567.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.
- Stephen E. Robertson and Hugo Zaragoza. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4):333–389, 2009. doi: 10.1561/15000000019.

- Alireza Salemi, Sheshera Mysore, Michael Bendersky, and Hamed Zamani. LaMP: When large language models meet personalization. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7370–7392, Bangkok, Thailand, 2024. Association for Computational Linguistics. arXiv:2304.11406.
- Ashudeep Singh and Thorsten Joachims. Fairness of exposure in rankings. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD '18)*, pages 2219–2228. ACM, 2018. doi: 10.1145/3219819.3220088.

A Supplementary Figure




Figure 7: Per-task detail underlying Finding 4: mean EU_{norm} difference (PL minus MMR) within shared ILD bins — a descriptive binned comparison, not a formal statistical matching procedure — by LaMP task and generator/retriever cell. Residuals are small and mostly negative, consistent with the pooled AME of Section 5.4; this per-task breakdown is supplementary detail, not an independent robustness claim.